

CS 5084 - Algorithms: Design and Analysis

Section 3.2 Graph Connectivity and Graph Traversal

Joe Johnson

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

Introduction

Now that we have some fundamental notions regarding graphs, we turn next to a basic algorithmic question: node-to-node connectivity.

Suppose we are given a graph $G = (V, E)$ and two particular nodes, s and t . We'd like to find an efficient algorithm that answers the following question:

Is there a path from s to t in G ? This is typically called the problem of determining s - t *connectivity*.

Introduction - contd.

There are two fundamental approaches to this problem:

- breadth-first search (BFS)
- depth-first search (DFS)

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

Breadth First Search

Breadth-first search (BFS) is the simplest algorithm for determining *s-t connectivity* based on the following idea:

In BFS, we search outward from our starting node, s , in all possible directions, adding nodes one layer at a time.

General BFS Algorithm:

- 1 Start with node s .
- 2 Visit each node v for which there exists an edge from s to v : *first layer*.
- 3 Visit each node w for which there exists an edge from a node in the first layer to w : *second layer*.
- 4 Repeat in this way until there are no more new, connected nodes.

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

BFS - Example

Consider the graph in diagram below:




Figure 3.2 In this graph, node 1 has paths to nodes 2 through 8, but not to nodes 9 through 13.

Let's run the BFS algorithm on this graph, starting with node, 1.

BFS - Example - contd.




Figure 3.2 In this graph, node 1 has paths to nodes 2 through 8, but not to nodes 9 through 13.

- 1 Start with node 1 as our starting node.
- 2 First layer: 2 and 3
- 3 Second layer: 4, 5, 7, 8
- 4 Third layer: 6
- 5 Terminate - no more new, connected nodes.

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - **BFS Layers**
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

Layers in BFS

We define the layers, L_1, L_2, L_3, \dots rigorously as follows:

- Layer L_0 consists of the starting node, s .
- Layer L_1 consists of all nodes that are neighbors of starting node, s .
- Assuming we have defined layers L_1, \dots, L_j then layer L_{j+1} consists of all nodes that do not belong to an earlier layer and that have an edge to a node in layer L_j .

Layers in BFS - contd.

(3.3) For each $j \geq 1$, layer L_j produced by BFS consists of all nodes at distance exactly j from s . There exists a path from node s to node t if and only if t appears in some layer.

Recall our definition of *distance* between nodes u and v which simply refers to the minimum number of edges (hops) on a path between u and v .

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

BFS Tree

BFS produces a tree T rooted at the starting node, s on the set of nodes reachable from s .

Consider the diagram below showing the construction of the BFS Tree for the example graph shown earlier:





Figure 3.2 In this graph, node 1 has paths to nodes 2 through 8, but not to nodes 9 through 13.



BFS Tree Theorem (3.4)

(3.4) Let T be a BFS Tree, let x and y be nodes in T belonging to layers L_i and L_j respectively, and let (x, y) be an edge e in graph, G . Then i and j differ by at most 1.

Proof: by Contradiction. Suppose i and j differ by more than 1, and in particular, $i < j - 1$. Now consider the point in the BFS algorithm when the edges incident to x are examined. Since x belongs to layer L_i , the only nodes discovered (or visited) by x are those nodes belonging to layers L_{i+1} and earlier. Since y is a neighbor of x , then y will be discovered by the time we get to layer L_{i+1} at the latest. But if $j = i + 1$, then that would imply that $i = j - 1$ (and thus, $i \not< j - 1$). Thus, we have a contradiction. Thus, i and j differ by at most 1.

Exploring a Connected Component

Definition: Consider the set of nodes that are reachable from starting node s . We call this set, R , the *connected component* of G containing s .


Note that BFS is just one approach that provides for a specific order in which we discover the nodes in R .

Also note that once we have the *connected component* for our starting node, s , we can simply find whether $t \in R$ in order to determine whether there exists a path from s to t .

At a more general level, we can build up the connected component, R , for some starting node s by “exploring” G in *any* order, starting from s .

- To start off we define $R = \{s\}$.
- Then, at any point in time if we find an edge (u, v) where $u \in R$ and $v \notin R$ we can add v to R .

Consider the diagram below:



Consider the following *general, nondeterministic* algorithm for finding the connected component, R :

```
 $R$  will consist of nodes to which  $s$  has a path  
Initially  $R = \{s\}$   
While there is an edge  $(u, v)$  where  $u \in R$  and  $v \notin R$   
    Add  $v$  to  $R$   
Endwhile
```

3.5 - The Algorithm for Building Connected Component, R

(3.5) The set R produced by this algorithm is precisely the connected component of G containing s .

Proof: We will show that both of the following are true:

- 1 For any node, v added to R by the algorithm, there exists a path from s to v .
- 2 For any node, w *not* added to R by the algorithm, there does *not* exist a path from s to w .

3.5 - Proof - contd.

- 1 Consider a situation in which there exists a path from s to u , and that there exists an edge from (u, v) . Thus, the algorithm will detect this edge, and add v to R . Does there exist a path from s to v , in this situation? Yes - namely, the path from s to u plus the edge from u to v .

3.5 - Proof - contd.

- ② by Contradiction. Suppose we have a node $w \notin R$ but that there *does* exist a path, P from s to w . Since $s \in R$ but $w \notin R$, there must be a first node v on P that does not belong to R . Assume this node is some node, v , where v is not equal to s . Thus, there is a node, u immediately preceding v on P , so (u, v) is an edge. Moreover, since v is the first node on P that does not belong to R , we must have $u \in R$. It follows that (u, v) is an edge where $u \in R$ and $v \notin R$. But this contradicts the stopping rule of the algorithm which will add any node v to R where there exists a node $u \in R$, and an edge (u, v) . In other words, the algorithm *will* add v to R , contrary to the assumption that $v \notin R$. Thus, if $w \notin R$, then, there *must* not exist a path, P , from s to w .

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

DFS Description

Consider a graph G as representing a maze of rooms, where each node represented a particular room and each edge represented a hallway leading from one room to another.

If you were wandering around in this maze, looking for a way to escape, you might start at a particular room, s , and try the first edge leading out of it, to a node, v . Then, you follow the first edge out of v and continue in this way until you reach a “dead end” - i.e., a node for which you had already explored all of its neighbors. At this point, you’d backtrack until you reached a node with an unexplored neighbor, and resume from there.

DFS Description - contd.

This approach to searching a graph is called the depth-first search algorithm (DFS).

Note that DFS, like BFS, is a particular implementation of the connected component algorithm for G containing s .

Recall that our original, motivating problem was to solve the s - t *connectivity problem*. We can apply $\text{DFS}(s)$, and find whether node t is in the *connectivity component* of G containing s using this algorithm.

DFS(u):

Mark u as "Explored" and add u to R

For each edge (u, v) incident to u

 If v is not marked "Explored" then

 Recursively invoke DFS(v)

 Endif

Endfor

Table of Contents

- 1 Introduction
- 2 Breadth First Search
 - Description
 - Example
 - BFS Layers
 - BFS Tree
- 3 Exploring a Connected Component
- 4 Depth First Search
 - Description
 - DFS Tree
- 5 The Set of All Connected Components

DFS Tree

Similar to the BFS algorithm, the DFS algorithm also produces a tree, T on the connectivity component of G containing s , rooted at s . However, it works a little differently:

- For a given node, u , we make u the parent of node v if u is responsible for discovering v . That is, if $\text{DFS}(u)$ results in a call to $\text{DFS}(v)$, then we create an edge in the DFS tree, T , from u to v , (u, v) .

DFS Tree - contd.

Recall our example graph, G :




Figure 3.2 In this graph, node 1 has paths to nodes 2 through 8, but not to nodes 9 through 13.

DFS Tree - contd.

Now, consider the construction of the DFS tree, based on the DFS algorithm:




(a)



(b)






(c)



(d)

DFS Tree - contd.



DFS Tree Observations

Consider the following observation about an DFS tree relative to a BFS tree:

- A BFS tree root-to-leaf paths are short and narrow while those for a DFS tree are long and deep. This reflects the nature of the respective algorithms.

(3.6) - Relationship between the DFS algorithm and the structure of the DFS tree

(3.6) For a given recursive call $\text{DFS}(u)$, all nodes that are marked “Explored” between the invocation and the end of this recursive call are descendants of u in T .

(3.7) - Structure of the DFS tree

(3.7) Let T be a DFS tree, let x and y be nodes in T , and let (x, y) be an edge of G that is *not* an edge of T . Then, one of x and y is an ancestor of the other.

(3.7) - Proof

Suppose that (x, y) is an edge of G that is not an edge of T and suppose without loss of generality that x is reached first by the DFS algorithm. When the edge (x, y) is examined during the execution of $\text{DFS}(x)$, it is not added to T because y is marked “Explored”. Since y was not marked “Explored” when $\text{DFS}(x)$ was first invoked, it is a node that was discovered between the invocation and end of the recursive call $\text{DFS}(x)$. It follows from (3.6) that y is a descendant of x .

The Set of All Connected Components

Until now, we've concerned ourselves with the connected component for a particular node, s . However, let's now consider the fact that all nodes in a graph, G , have a connected component.

What is the relationship between the connected components among the various nodes in a graph?

(3.8) - A highly structured relationship between connected components in a graph

(3.8) For any two nodes s and t in a graph, their connected components are either *identical* or *disjoint*.

(3.8) - Proof

Proof: We will complete this proof in two steps:

- 1 First, we'll show that if there exists a path between any two nodes s and t , then the connected components of s and t are *identical*.
- 2 Second, we'll show that if there does *not* exist a path between two nodes s and t , then their connected components are *disjoint*.

(3.8) - Proof (contd.)

- 1 Suppose there exists a path between any two nodes s and t . Now, let's suppose there exists a path between s and some other node, v , (and thus, v is a member of the connected component for s). Question: Is there a path from t to v ? Answer: Yes, because we can make a path that consists first of the path from t to s , and then from s to v . So, this implies that v must also be in the connected component of t . So, for any node v in the connected component of s , v is also in the connected component of t , and vice versa by the same reasoning.

(3.8) - Proof (contd.)

- 1 Suppose there does *not* exist a path between two nodes s and t . Now, suppose, however, there exists a node v such that there exists a path from s to v and a path from t to v . Under this assumption of such a node, v , there *does* exist a path from s to t - namely, a path leading from s to v and then from v to t . Since our original assumption is that there is *no* path between s and t , such a node, v , cannot exist. Thus, if there is no path between two nodes, then they share no nodes in common in their respective connected components. Thus, the connected components for s and t are disjoint.